

Semantic fashion recommendation service: data flow

natural-language request -> LLM plan -> hybrid retrieval per outfit slot -> LLM rerank -> ranked products

1. Offline: build the index once (make data / ingest / index)



Deployments (Railway): the container downloads this tarball at boot when INDEX_URL + INDEX_SHA256 are set; checksum, safe extraction, atomic install.

2. Online: one request (POST /recommend), one deadline (40 s)

